

Stage	Description	Model Source	Prompt Logic	Outputs	Visibility			
R1	Initial draft round; all selected LLMs respond to user query	ACTIVE models selected from READY	Each ACTIVE model receives the same query and generates an independent draft.	INITIAL outputs	User-visible AFTER ULTRA (tab: INITIAL)			
R2	META revision round; each model receives peer responses to reconsider its own answer	Same ACTIVE models as R1	Instruct: 'Do not assume any response is true. Revise your answer after reviewing your peers. List contradictions resolved and what you changed.'	META outputs	User-visible AFTER ULTRA (tab: META)			
R3	Final synthesis round; neutral model combines META drafts into unified synthesis	ULTRA model selected from READY by preference	Instruct: 'Review all META drafts. Merge convergent points, resolve contradictions, and cite which META claims were retained or omitted. Generate one coherent synthesis with confidence intervals and model statistics.'	ULTRA output + STATS	User-visible (tab: ULTRA)			
STATS	Collect per-run analytics (counts	timings	agreement/revision rates	errors)	System aggregation	Automatically derived from stored run artifacts.	stats.json	User-visible (tab: STATS)
Error Handling	Handle missing/failed models and degraded runs	Engine auto-handles READY<2 or ULTRA missing	If READY < 2: abort with 'low pluralism' warning. If ULTRA missing: degraded mode flag.	status.json, error logs	User-visible (summary banner)			